

wrangle_act

October 12, 2018

1 Project: “We Rate Dogs” Data Analysis

1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

Reference

Introduction

Twitter account ‘WeRateDogs (@dog_rates)’ is famous for its unique tweets which evaluate dog pictures which are posted from readers. In this analysis, we get the tweets data from Twitter API and the additional data about pictures of those tweets and assess them. Specifically, we focus on 2 topics below

- How the ratings are related with stages and types of dogs.
- Relationship between favorite counts and the ratings.

Data Wrangling

1.2 Gather

```
In [188]: #import packages
```

```
import pandas as pd
import numpy as np
import requests
import tweepy
import json
import timeit
import datetime
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
%matplotlib inline
```

1.2.1 Acquire Twitter archive data

```
In [189]: # read "twitter-archive-enhanced.csv" file
df_archive = pd.read_csv("twitter-archive-enhanced.csv")
df_archive.columns

Out[189]: Index(['tweet_id', 'in_reply_to_status_id', 'in_reply_to_user_id', 'timestamp',
                'source', 'text', 'retweeted_status_id', 'retweeted_status_user_id',
                'retweeted_status_timestamp', 'expanded_urls', 'rating_numerator',
                'rating_denominator', 'name', 'doggo', 'floofer', 'pupper', 'puppo'],
                dtype='object')
```

1.2.2 Acquire image-predictions data

```
In [190]: # download image prediction tsv file as "image_predictions.tsv" in the local directory
url = "https://d17h27t6h515a5.cloudfront.net/topher/2017/August/599fd2ad_image-predictions.tsv"
response = requests.get(url)

with open("image_predictions.tsv", mode="wb") as file:
    file.write(response.content)

In [191]: # import the file into this notebook as a dataframe
df_image = pd.read_csv("image_predictions.tsv", sep="\t")
df_image.columns

Out[191]: Index(['tweet_id', 'jpg_url', 'img_num', 'p1', 'p1_conf', 'p1_dog', 'p2',
                'p2_conf', 'p2_dog', 'p3', 'p3_conf', 'p3_dog'],
                dtype='object')
```

1.2.3 Acquire tweet json-file (from TwitterAPI)

```
In [5]: # get api using tweepy
consumer_key = "####"
consumer_secret = "####"
access_token = "####"
access_token_secret = "####"

auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
auth.set_access_token(access_token, access_token_secret)
api = tweepy.API(auth)

In [11]: # select tweets we need in api json and get "retweet_count" and "favorite_count" info
# store them in the list "data_list"
start = timeit.timeit()

data_list = []
for tweet_id in df_archive["tweet_id"]:
    try:
        api_status = api.get_status(tweet_id, tweet_mode='extended', wait_on_rate_limit=True)
```

```

        tweet_data = {}
        tweet_data["tweet_id"] = api_status._json["id"]
        tweet_data["ret"] = api_status._json["retweet_count"]
        tweet_data["fav"] = api_status._json["favorite_count"]
        data_list.append(tweet_data)
    except Exception as e:
        print(str(tweet_id) + " : " + str(e))

end = timeit.timeit()
print(end - start)

888202515573088257 : [{'code': 144, 'message': 'No status found with that ID.'}]
873697596434513921 : [{'code': 144, 'message': 'No status found with that ID.'}]
872668790621863937 : [{'code': 144, 'message': 'No status found with that ID.'}]
869988702071779329 : [{'code': 144, 'message': 'No status found with that ID.'}]
866816280283807744 : [{'code': 144, 'message': 'No status found with that ID.'}]

```

Rate limit reached. Sleeping for: 648

```

861769973181624320 : [{'code': 144, 'message': 'No status found with that ID.'}]
845459076796616705 : [{'code': 144, 'message': 'No status found with that ID.'}]
842892208864923648 : [{'code': 144, 'message': 'No status found with that ID.'}]
837012587749474308 : [{'code': 144, 'message': 'No status found with that ID.'}]
827228250799742977 : [{'code': 144, 'message': 'No status found with that ID.'}]
802247111496568832 : [{'code': 144, 'message': 'No status found with that ID.'}]
775096608509886464 : [{'code': 144, 'message': 'No status found with that ID.'}]
771004394259247104 : [{'code': 179, 'message': 'Sorry, you are not authorized to see this status.'}]
770743923962707968 : [{'code': 144, 'message': 'No status found with that ID.'}]
754011816964026368 : [{'code': 144, 'message': 'No status found with that ID.'}]

```

Rate limit reached. Sleeping for: 647

Rate limit reached. Sleeping for: 563

0.002290688003995456

```

In [59]: # write the list in 'tweet_json.txt' (each element of tweet on each line in the file)
        with open('tweet_json.txt', 'w') as file:
            for line in data_list:
                json.dump(line, file)
                file.write("\n")

In [192]: # load 'tweet_json.txt' and store it in the list "json_list" (each element of this list
        # this list should be the same as the list "data_list"
        json_list = []

```

```

with open('tweet_json.txt','r') as file:
    for line in file:
        json_list.append(json.loads(line))

```

```

In [193]: # construct dataframe from the list "json_list" and sort columns
df_api = pd.DataFrame(json_list)
df_api = df_api.iloc[:,[1,0,2]]
df_api.columns

```

```

Out[193]: Index(['ret', 'fav', 'tweet_id'], dtype='object')

```

1.3 Assess

```

In [194]: # visual assessment
df_archive.head(5)

```

```

Out[194]:
      tweet_id  in_reply_to_status_id  in_reply_to_user_id \
0  892420643555336193                NaN                NaN
1  892177421306343426                NaN                NaN
2  891815181378084864                NaN                NaN
3  891689557279858688                NaN                NaN
4  891327558926688256                NaN                NaN

```

```

      timestamp \
0  2017-08-01 16:23:56 +0000
1  2017-08-01 00:17:27 +0000
2  2017-07-31 00:18:03 +0000
3  2017-07-30 15:58:51 +0000
4  2017-07-29 16:00:24 +0000

```

```

      source
0  <a href="http://twitter.com/download/iphone" rel="nofollow">Twitter for iPhone</a>
1  <a href="http://twitter.com/download/iphone" rel="nofollow">Twitter for iPhone</a>
2  <a href="http://twitter.com/download/iphone" rel="nofollow">Twitter for iPhone</a>
3  <a href="http://twitter.com/download/iphone" rel="nofollow">Twitter for iPhone</a>
4  <a href="http://twitter.com/download/iphone" rel="nofollow">Twitter for iPhone</a>

```

```

0          This is Phineas. He's a mystical boy. Only
1  This is Tilly. She's just checking pup on you. Hopes you're doing ok. If not, she
2      This is Archie. He is a rare Norwegian Pouncing Corgo. Lives in the tall grass
3          This is Darla. She commenced a snooz
4  This is Franklin. He would like you to stop calling him "cute." He is a very fier

```

```

      retweeted_status_id  retweeted_status_user_id  retweeted_status_timestamp \
0                NaN                NaN                NaN
1                NaN                NaN                NaN
2                NaN                NaN                NaN
3                NaN                NaN                NaN

```

4	NaN	NaN	NaN
---	-----	-----	-----

0		https://twitter.co
1		https://twitter.co
2		https://twitter.co
3		https://twitter.co
4	https://twitter.com/dog_rates/status/891327558926688256/photo/1,	https://twitter.co

	rating_numerator	rating_denominator	name	doggo	floofer	pupper	puppo
0	13	10	Phineas	None	None	None	None
1	13	10	Tilly	None	None	None	None
2	12	10	Archie	None	None	None	None
3	13	10	Darla	None	None	None	None
4	12	10	Franklin	None	None	None	None

In [195]: df_archive.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id    78 non-null float64
in_reply_to_user_id      78 non-null float64
timestamp               2356 non-null object
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id      181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls            2297 non-null object
rating_numerator         2356 non-null int64
rating_denominator       2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
dtypes: float64(4), int64(3), object(10)
memory usage: 313.0+ KB
```

In [196]: # visual assessment
df_api.sample(5)

```
Out[196]:
```

	ret	fav	tweet_id
366	184	2336	828361771580813312
459	5205	0	816829038950027264
2072	246	752	670797304698376195

```

905    3941    11093    756526248105566208
2015    3624     7237    671789708968640512

```

```
In [197]: df_api.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2341 entries, 0 to 2340
Data columns (total 3 columns):
ret          2341 non-null int64
fav          2341 non-null int64
tweet_id     2341 non-null int64
dtypes: int64(3)
memory usage: 54.9 KB

```

```
In [198]: # visual assessment
          df_image.sample(5)
```

```

Out[198]:
          tweet_id          jpg_url \
1819  834209720923721728  https://pbs.twimg.com/media/C501UAaWIAAMBd.jpg
2024  881666595344535552  https://pbs.twimg.com/media/DDxPFwbWAAEbVVR.jpg
2020  880872448815771648  https://pbs.twimg.com/media/DDl8zzJW0AAisCJ.jpg
472   675145476954566656  https://pbs.twimg.com/media/CV6ZOPqWsAA20Uj.jpg
1320  756288534030475264  https://pbs.twimg.com/media/Cn7gaHrWIAAZJMt.jpg

          img_num          p1  p1_conf  p1_dog          p2 \
1819          1  golden_retriever  0.754799    True    Pekinese
2024          1          Saluki  0.529012    True  Afghan_hound
2020          1    Pembroke  0.791416    True  Norwich_terrier
472          1  Labrador_retriever  0.458746    True    Great_Dane
1320          3          conch  0.925621   False  French_bulldog

          p2_conf  p2_dog          p3  p3_conf  p3_dog
1819  0.197861    True  Labrador_retriever  0.008654    True
2024  0.250003    True    golden_retriever  0.160739    True
2020  0.061393    True    Chihuahua  0.033726    True
472   0.235504    True  Staffordshire_bullterrier  0.116864    True
1320  0.032492    True    tiger_cat  0.006679   False

```

```
In [199]: df_image.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id     2075 non-null int64
jpg_url      2075 non-null object
img_num      2075 non-null int64
p1           2075 non-null object
p1_conf      2075 non-null float64

```

```

p1_dog      2075 non-null bool
p2          2075 non-null object
p2_conf     2075 non-null float64
p2_dog      2075 non-null bool
p3          2075 non-null object
p3_conf     2075 non-null float64
p3_dog      2075 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB

```

1.3.1 Quality

df_archive table

- “timestamp” is a string not a datetime.
- Rows of tweets which are later than 08/01 2017 should be removed.
- Tweets which are original ratings should be extracted.
- “name” column are unreliable, thus it should be delete.
- ‘rating_numerator’ and ‘rating_denominator’ columns are not necessarily correctly extracted.
- ‘rating’ column which represents (‘rating_numerator’ / rating_denominator) should be created.
- Stage columns (‘doggo’, ‘floofer’, ‘pupper’, ‘puppo’) are not necessarily correctly extracted.

df_image table

- Some pictures are predicted not as dogs. If a picture is not predicted as dog till the 3rd prediction, delete that row.

1.3.2 Tidiness

- “df_api” and “df_image” should be merged to “df_archive”
- Dog stages in the archive data should be in 1 column.
- We need only the most primary confident prediction of dog types from pictures, so make the column “predicted dog type” in place of p1~p3 predictions.

1.4 Clean

```

In [200]: # copy the dataframes for the following cleanness
          archive_clean = df_archive.copy()
          image_clean = df_image.copy()
          api_clean = df_api.copy()

```

1.4.1 Tweets which are original should be extracted from archive table.

```

In [201]: # use .query method to extract rows whose both "in_reply_to_status_id" and "retweeted"
          archive_clean = archive_clean[archive_clean.in_reply_to_status_id.isnull() & archive_clean.retweeted.isnull()]
          archive_clean.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2097 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2097 non-null int64
in_reply_to_status_id   0 non-null float64
in_reply_to_user_id     0 non-null float64
timestamp               2097 non-null object
source                  2097 non-null object
text                    2097 non-null object
retweeted_status_id     0 non-null float64
retweeted_status_user_id 0 non-null float64
retweeted_status_timestamp 0 non-null object
expanded_urls           2094 non-null object
rating_numerator        2097 non-null int64
rating_denominator      2097 non-null int64
name                    2097 non-null object
doggo                   2097 non-null object
floofer                 2097 non-null object
pupper                 2097 non-null object
puppo                   2097 non-null object
dtypes: float64(4), int64(3), object(10)
memory usage: 294.9+ KB

```

```

In [202]: # remove columns related to "in_reply_status" or "retweeted_status" which is no longer needed
archive_clean = archive_clean.drop(columns=["in_reply_to_status_id", "in_reply_to_user_id",
                                             "retweeted_status_id", "retweeted_status_user_id", "retweeted_status_timestamp"])

```

```

In [203]: # check columns
archive_clean.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2097 entries, 0 to 2355
Data columns (total 12 columns):
tweet_id                2097 non-null int64
timestamp               2097 non-null object
source                  2097 non-null object
text                    2097 non-null object
expanded_urls           2094 non-null object
rating_numerator        2097 non-null int64
rating_denominator      2097 non-null int64
name                    2097 non-null object
doggo                   2097 non-null object
floofer                 2097 non-null object
pupper                 2097 non-null object
puppo                   2097 non-null object
dtypes: int64(3), object(9)
memory usage: 213.0+ KB

```



```
In [204]: archive_clean.columns
```

```
Out[204]: Index(['tweet_id', 'timestamp', 'source', 'text', 'expanded_urls',  
               'rating_numerator', 'rating_denominator', 'name', 'doggo', 'floofer',  
               'pupper', 'puppo'],  
              dtype='object')
```

1.4.2 “name” column in archive data is not necessarily in this analysis so delete this.

```
In [205]: # use drop() method  
archive_clean = archive_clean.drop(columns="name")
```

```
In [206]: # confirm that there is no column named "name" now.  
assert "name" not in archive_clean.columns
```

1.4.3 “timestamp” in archive data is a string not a datetime.

```
In [207]: # use pd.to_datetime() function to convert the data type.  
archive_clean.timestamp = pd.to_datetime(archive_clean.timestamp)
```

```
In [208]: # confirm it was converted properly and the term in it is possible  
archive_clean.timestamp.describe()
```

```
Out[208]: count                2097  
unique                2097  
top      2016-09-12 15:10:21  
freq                      1  
first    2015-11-15 22:32:08  
last     2017-08-01 16:23:56  
Name: timestamp, dtype: object
```

1.4.4 Rows of tweets which are later than 08/01 2017 should be removed from archive data.

```
In [209]: # use query() method to extract rows  
archive_clean = archive_clean.query('timestamp < datetime.date(2017, 8, 1)')
```

```
In [210]: # confirm the latest tweet information in archive data is earlier than 08/01 2017  
archive_clean.timestamp.describe()
```

```
Out[210]: count                2095  
unique                2095  
top      2016-09-12 15:10:21  
freq                      1  
first    2015-11-15 22:32:08  
last     2017-07-31 00:18:03  
Name: timestamp, dtype: object
```

1.4.5 'rating_numerator' and 'rating_denominator' columns in archive data are not necessarily correctly extracted.

Definition

We take a look at "text" column in the archive data. We can assume that ratings are in the form of fractions, so extract the fractions as ratings.

However, we will assess texts of tweets visually in the case where the texts have 2 or more fractions.

When we find that there are 2 or more different dogs' ratings in 1 tweet in this process, delete the tweet because it can be misleading.

```
In [211]: # Firstly, extract the first fraction which appears in the "text" and confirm that w
fractions = archive_clean.text.str.extract(r'([0-9]+[.]*[0-9]*[/][0-9]+)')
fractions.isnull().sum()
```

```
Out[211]: 0      0
dtype: int64
```

```
In [212]: # split the fractions into numerators and denominators
f_num, f_deno = fractions.iloc[:,0].str.split("/").str
archive_clean.rating_numerator = f_num.astype(float)
archive_clean.rating_denominator = f_deno.astype(float)
archive_clean.rating_denominator.dtype
```

```
Out[212]: dtype('float64')
```

```
In [213]: # extract fractions from "text" column using .str.extractall().
# If there are 2 or more than "text" column, assess each row one by one.
pd.options.display.max_colwidth = 150

ratings = archive_clean.text.str.extractall(r'([0-9]+[.]*[0-9]*[/][0-9]+)')

list_duplicated = [] # list of index whose rows have 2 or more fractions the "te
for multi_index in ratings.index:
    if multi_index[1] != 0:
        list_duplicated.append(multi_index[0])

print(list_duplicated)
archive_clean.info()
```

```
[766, 1007, 1068, 1165, 1202, 1222, 1359, 1459, 1465, 1508, 1525, 1538, 1662, 1795, 1832, 1897
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2095 entries, 2 to 2355
Data columns (total 11 columns):
tweet_id      2095 non-null int64
timestamp     2095 non-null datetime64[ns]
source        2095 non-null object
text          2095 non-null object
```

```

expanded_urls      2092 non-null object
rating_numerator   2095 non-null float64
rating_denominator 2095 non-null float64
doggo              2095 non-null object
floofer            2095 non-null object
pupper            2095 non-null object
puppo              2095 non-null object
dtypes: datetime64[ns](1), float64(2), int64(1), object(7)
memory usage: 196.4+ KB

```

```

In [214]: # assess each row one by one
          archive_clean.text[list_duplicated]

```

```

Out[214]: 766                                     "Yep... just as I suspected. You're not floss
1007                                     This is Bookstore and Seaweed. Bookstore is tired and Sea
1068                                     After so many requests, this is Bretagne. She was the last surviving
1165
1202                                     This is Bluebert. He just saw that both #Final
1222                                     Meet Travis and Flurp. Travis is pretty chill but Flurp can't lie down pro
1359                                     This is Socks. That water pup w the super legs just splashed him
1459                                     This may be the greatest video I've ever been sent. 4/10 for Charles
1465                                     Meet Oliviér. He takes killer selfies. Has a dog of his own. It leaps at ran
1508                                     When bae says they can't go out but you see them with someone else that s
1525                                     This is Eriq. His friend just reminded him of last year's super bowl. No
1538                                     Meet Fynn & Taco. Fynn is an all-powerful leaf lord and Taco is in the w
1662                                     This is Darrel. He just robbed a 7/11 and is in a high speed police
1795                                     Meet Tassy & Bee. Tassy is pretty chill, but Bee is convinced the Ruffi
1832                                     These two pups just met and have instantly bonded. Spectacular scen
1897                                     Meet Rufio. He is unaware of the pink legless pupper wrapped around him. M
1901                                     Two gorgeous dogs here. Little waddling dog is a rebel. Refuses to loo
1970                                     Meet Eve. She's a raging alcoholic 8/10 (would b 11/10 but pupper alco
2010                                     10/10 for dog. 7/10 for
2010                                     10/10 for dog. 7/10 for
2064                                     Meet Holly. She's trying to teach small human-like pup about blocks but l
2113                                     Meet Hank and Sully. Hank is very proud of the pumpkin they found
2177                                     Here we have Pancho and Peaches. Pancho is a Condoleezza Gryffindor, an
2216                                     This is Spark. He's nervous. Other dog hasn't moved in a while. Won't come
2263                                     This is Kial. Kial is either wearing a cape, which would be rad, or
2272                                     Two dogs in this one. Both are rare Jujitsu Pythagoreans. One slight
2306                                     These are Peruvian Feldspars. Their names are Cupit and Prencer. Both res
2335                                     This is an Albanian 3 1/2 legged Episcopalian. Loves well-polish
Name: text, dtype: object

```

```

In [215]: def put_ratings(index, numerator, denominator):
           archive_clean.ix[index, "rating_denominator"] = denominator
           archive_clean.ix[index, "rating_numerator"] = numerator

```

```
indexes_to_be_deleted = [766, 1007, 1165, 1222, 1459, 1525,1538,1795,1832,1901,2113,
archive_clean = archive_clean.drop(index = indexes_to_be_deleted)
```

```
put_ratings(1068,14,10)
put_ratings(1202,11,10)
put_ratings(1359,9,10)
put_ratings(1465,10,10)
put_ratings(1508,10,10)
put_ratings(1662,10,10)
put_ratings(1897,10,10)
put_ratings(1970,8,10)
put_ratings(2010,10,10)
put_ratings(2064,11,10)
put_ratings(2263,10,10)
put_ratings(2335, 9,10)
```

```
archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2080 entries, 2 to 2355
Data columns (total 11 columns):
tweet_id          2080 non-null int64
timestamp         2080 non-null datetime64[ns]
source            2080 non-null object
text              2080 non-null object
expanded_urls     2077 non-null object
rating_numerator  2080 non-null float64
rating_denominator 2080 non-null float64
doggo             2080 non-null object
floofer           2080 non-null object
pupper            2080 non-null object
puppo             2080 non-null object
dtypes: datetime64[ns](1), float64(2), int64(1), object(7)
memory usage: 275.0+ KB
```

```
/anaconda3/lib/python3.6/site-packages/ipykernel/__main__.py:2: DeprecationWarning:
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing
```

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>

```
from ipykernel import kernelapp as app
```

```
/anaconda3/lib/python3.6/site-packages/ipykernel/__main__.py:3: DeprecationWarning:
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing
```

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>
`app.launch_new_instance()`

1.4.6 'rating' column which represents (rating_numerator / rating_denominator) should be created.

```
In [216]: # create a new column called "rating" and drop 'rating_numerator' and 'rating_denominator'
archive_clean["rating"] = archive_clean.rating_numerator / archive_clean.rating_denominator
archive_clean = archive_clean.drop(columns=['rating_numerator', 'rating_denominator'])
archive_clean.columns
```

```
Out[216]: Index(['tweet_id', 'timestamp', 'source', 'text', 'expanded_urls', 'doggo',
                'floofer', 'pupper', 'puppo', 'rating'],
                dtype='object')
```

```
In [217]: # Also, I will check "rating" extremely high or low and remove inappropriate ones.
archive_clean[(archive_clean.rating < 0.1) | (archive_clean.rating > 2.0)].text
```

```
Out[217]: 315                                When you're so blinded by your systematic
516      Meet Sam. She smiles 24/7 & secretly aspires to be a reindeer. \nKeep Sam
979
2074                                This :
                                     After so m
Name: text, dtype: object
```

```
In [218]: # We delete tweet whose picture are not dogs.
archive_clean = archive_clean.drop([315, 516, 2074])
archive_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2077 entries, 2 to 2355
Data columns (total 10 columns):
tweet_id      2077 non-null int64
timestamp     2077 non-null datetime64[ns]
source        2077 non-null object
text          2077 non-null object
expanded_urls 2074 non-null object
doggo         2077 non-null object
floofer       2077 non-null object
pupper        2077 non-null object
puppo         2077 non-null object
rating        2077 non-null float64
dtypes: datetime64[ns](1), float64(1), int64(1), object(7)
memory usage: 178.5+ KB
```

1.4.7 Stage columns ('doggo', 'floofer', 'pupper', 'puppo') in archive are not necessarily correctly extracted.

```
In [219]: # First, delete these columns
```

```
archive_clean = archive_clean.drop(columns = ['doggo', 'floofer', 'pupper', 'puppo'])
archive_clean.columns
```

```
Out[219]: Index(['tweet_id', 'timestamp', 'source', 'text', 'expanded_urls', 'rating'], dtype=
```

```
In [220]: # add the new column of each stage (boolean)
```

```
archive_clean["doggo"] = archive_clean.text.str.contains("doggo")
archive_clean["pupper"] = archive_clean.text.str.contains("pupper")
archive_clean["puppo"] = archive_clean.text.str.contains("puppo")
archive_clean["floof"] = archive_clean.text.str.contains("floof")
archive_clean["snoot"] = archive_clean.text.str.contains("snoot")
archive_clean["blep"] = archive_clean.text.str.contains("blep")

archive_clean[['doggo', 'pupper', 'puppo', 'floof', 'snoot', 'blep']].sum()
```

```
Out[220]: doggo      84
pupper    245
puppo     29
floof     23
snoot      0
blep       1
dtype: int64
```

```
In [221]: # assess the rows which has 2 or more stage name in the "text" column.
```

```
archive_clean.text[archive_clean[['doggo', 'pupper', 'puppo', 'floof', 'snoot', 'blep']]
```

```
Out[221]: 172      I have stumbled puppon a doggo painting party. They're looking to be the next
191      Here's a puppo participating in the #ScienceMarch. Cleverly disguising her o
531      Here we have Burke (pupper) and Dexter (doggo). Pupper wants to be exact
575      This is Bones. He's being haunted by another doggo of roughly the same s
705      This is Pinot. He's a sophisticated doggo. You can tell by the hat. Also poin
889      Meet Maggie & Lila. Maggie is the doggo, Lila is the pupper. They ar
956      Please stop sending it pictures that don't even have a doggo or
1063      This is just downr
1113      L.
Name: text, dtype: object
```

```
In [222]: # read each text and choose the most appropriate stage.
```

```
# If we find that there are 2 dogs in the picture and thus there are 2 stages, delet
# This is because these rows can mis leading when we consider them with image predic
```

```
def cancel_stage(index, stage):
    archive_clean.ix[index, stage] = False
```

```

indexes_to_be_deleted2 = [531, 889, 1062, 1113]
archive_clean = archive_clean.drop(index = indexes_to_be_deleted2)

cancel_stage(172, "puppo")
cancel_stage(191, "doggo")
cancel_stage(575, "doggo")
cancel_stage(705, "pupper")
cancel_stage(956, "doggo")
cancel_stage(1063, "pupper")

```

/anaconda3/lib/python3.6/site-packages/ipykernel/__main__.py:6: DeprecationWarning:
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>

In [223]: *# confirm that there is no row which has 2 or more stages*

```

assert len(archive_clean[archive_clean[['doggo', 'pupper', 'puppo', 'floof', 'snoot', '']

```

1.4.8 Dog stages in the archive data should be in 1 column.

First, we separate the rows into those which is given stages(df“have_stage_tweets”) if dog and those which are not(df“no_stage_tweets”).

- For df“no_stage_tweets”, we give a “stage” column which stores “NaN”.
- For df“have_stage_tweets”, we give a “stage” column which stores the name of stage.

Finally, we concatbnate these dataframes to get the final dataframe.

In [224]: *# extract the rows which are not assigns dos stages as "no_stage_tweets".*

```

# construct new column named "stage" which is stored NaN and delete unnecessary columns

```

```

no_stage_tweets = archive_clean[archive_clean[['doggo', 'pupper', 'puppo', 'floof', 'snoot', '']]
no_stage_tweets = no_stage_tweets.drop(columns = ['doggo', 'pupper', 'puppo', 'floof', 'snoot', ''])
no_stage_tweets["stage"] = np.nan
no_stage_tweets.columns

```

Out[224]: Index(['tweet_id', 'timestamp', 'source', 'text', 'expanded_urls', 'rating',
'stage'],
dtype='object')

In [225]: *# extract the rows which are assigns any dos stages as "have_stage_tweets".*

```

# construct new column named "stage" and delete unnecessary columns (using .melt())

```

```

have_stage_tweets = archive_clean[archive_clean[['doggo', 'pupper', 'puppo', 'floof', '']]

## confirm that all rows are in df"no_stage_tweets" or df"have_stage_tweets".
assert len(no_stage_tweets) + len(have_stage_tweets) == len(archive_clean)

have_stage_tweets = have_stage_tweets.melt(id_vars=['tweet_id', 'timestamp', 'source'])
have_stage_tweets = have_stage_tweets[have_stage_tweets.value == True]
have_stage_tweets = have_stage_tweets.drop(columns = 'value')
have_stage_tweets

## confirm that df"no_stage_tweets" or df"have_stage_tweets" and same columns.
no_stage_tweets.columns == have_stage_tweets.columns

```

```
Out[225]: array([ True,  True,  True,  True,  True,  True,  True])
```

```

In [226]: # concatenate df"no_stage_tweets" or df"have_stage_tweets"
archive_clean2 = pd.concat([no_stage_tweets, have_stage_tweets]).sort_index().reset_index()
assert len(archive_clean2) == len(archive_clean)

pd.options.display.max_colwidth = 30
archive_clean2.head()

## reset index
archive_clean2 = archive_clean2.reset_index(drop=True)
archive_clean2.head()

```

```

Out[226]:
      tweet_id      timestamp      source \
0  890240255349198849  2017-07-26 15:59:51  <a href="http://twitter.co...
1  891815181378084864  2017-07-31 00:18:03  <a href="http://twitter.co...
2  891689557279858688  2017-07-30 15:58:51  <a href="http://twitter.co...
3  891327558926688256  2017-07-29 16:00:24  <a href="http://twitter.co...
4  884162670584377345  2017-07-09 21:29:42  <a href="http://twitter.co...

      text      expanded_urls  rating  stage
0  This is Cassie. She is a c...  https://twitter.com/dog_ra...    1.4  doggo
1  This is Archie. He is a ra...  https://twitter.com/dog_ra...    1.2   NaN
2  This is Darla. She commenc...  https://twitter.com/dog_ra...    1.3   NaN
3  This is Franklin. He would...  https://twitter.com/dog_ra...    1.2   NaN
4  Meet Yogi. He doesn't have...  https://twitter.com/dog_ra...    1.2  doggo

```

1.4.9 Some pictures are predicted not as dogs. If a picture is not predicted as dog till the 3rd prediction, delete that row in prediction image data.

```

In [227]: # use .query() method to extract rows of tweets which are predicted as dogs
image_clean = image_clean.query(' p1_dog == True | p2_dog == True | p3_dog == True ')

In [228]: # confirm there is no row which is not predicted as dogs either in p1-p3 predictions
assert ( (image_clean.p1_dog | image_clean.p2_dog | image_clean.p3_dog)==0 ).sum() == 0
image_clean.info()

```



```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1751 entries, 0 to 2073
Data columns (total 12 columns):
tweet_id      1751 non-null int64
jpg_url       1751 non-null object
img_num       1751 non-null int64
p1            1751 non-null object
p1_conf       1751 non-null float64
p1_dog        1751 non-null bool
p2            1751 non-null object
p2_conf       1751 non-null float64
p2_dog        1751 non-null bool
p3            1751 non-null object
p3_conf       1751 non-null float64
p3_dog        1751 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 141.9+ KB

```

1.4.10 We need only the most primary confident prediction of dog types from pictures, so make the column “predicted dog type” in place of p1~p3 predictions.

```

In [229]: # create a new list of dog types according to p1 ~p3.
          # add new column "dog_type" in df"image_clean" using it.

```

```

image_clean = image_clean.reset_index(drop=True )
dog_type_list = []

for index in image_clean.index:
    if image_clean.p1_dog.iloc[index] == True:
        dog_type_list.append(image_clean.p1.iloc[index])
    elif image_clean.p2_dog.iloc[index] == True:
        dog_type_list.append(image_clean.p2.iloc[index])
    elif image_clean.p3_dog.iloc[index] == True:
        dog_type_list.append(image_clean.p3.iloc[index])
    else:
        raise # throw an error when there is no dog type predicted in p1~p3

```

```

image_clean["dog_type"] = dog_type_list

```

```

In [230]: # remove p1~p3 columns because the are no longer needed.

```

```

image_clean = image_clean.drop(columns = ["p1", "p1_conf", "p1_dog", "p2", "p2_conf", "p2_dog", "p3", "p3_conf", "p3_dog"])

```

```

In [231]: # confirm the columns

```

```

image_clean.head()

```

```
Out [231]:
```

	tweet_id	jpg_url	img_num	\
0	666020888022790149	https://pbs.twimg.com/medi...	1	
1	666029285002620928	https://pbs.twimg.com/medi...	1	
2	666033412701032449	https://pbs.twimg.com/medi...	1	
3	666044226329800704	https://pbs.twimg.com/medi...	1	
4	666049248165822465	https://pbs.twimg.com/medi...	1	

	dog_type
0	Welsh_springer_spaniel
1	redbone
2	German_shepherd
3	Rhodesian_ridgeback
4	miniature_pinscher

1.4.11 "df_api" and "df_image" should be merged to "df_archive"

```
In [232]: # use inner merge to merge 3 dataframes to make final dataframe "df"
```

```
df = archive_clean2.merge(api_clean)
df = df.merge(image_clean)
df.head()
```

```
Out [232]:
```

	tweet_id	timestamp	source	\
0	890240255349198849	2017-07-26 15:59:51	<a href="http://twitter.co...	
1	891815181378084864	2017-07-31 00:18:03	<a href="http://twitter.co...	
2	891689557279858688	2017-07-30 15:58:51	<a href="http://twitter.co...	
3	891327558926688256	2017-07-29 16:00:24	<a href="http://twitter.co...	
4	884162670584377345	2017-07-09 21:29:42	<a href="http://twitter.co...	

	text	expanded_urls	rating	\
0	This is Cassie. She is a c...	https://twitter.com/dog_ra...	1.4	
1	This is Archie. He is a ra...	https://twitter.com/dog_ra...	1.2	
2	This is Darla. She commenc...	https://twitter.com/dog_ra...	1.3	
3	This is Franklin. He would...	https://twitter.com/dog_ra...	1.2	
4	Meet Yogi. He doesn't have...	https://twitter.com/dog_ra...	1.2	

	stage	ret	fav	jpg_url	img_num	\
0	doggo	7300	31555	https://pbs.twimg.com/medi...	1	
1	NaN	4102	24746	https://pbs.twimg.com/medi...	1	
2	NaN	8540	41677	https://pbs.twimg.com/medi...	1	
3	NaN	9265	39842	https://pbs.twimg.com/medi...	2	
4	doggo	2952	20137	https://pbs.twimg.com/medi...	1	

	dog_type
0	Pembroke
1	Chihuahua
2	Labrador_retriever
3	basset

4 German_shepherd

```
In [233]: # confirm final dataframe
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1645 entries, 0 to 1644
Data columns (total 12 columns):
tweet_id      1645 non-null int64
timestamp     1645 non-null datetime64[ns]
source        1645 non-null object
text          1645 non-null object
expanded_urls 1645 non-null object
rating        1645 non-null float64
stage         287 non-null object
ret           1645 non-null int64
fav           1645 non-null int64
jpg_url       1645 non-null object
img_num       1645 non-null int64
dog_type      1645 non-null object
dtypes: datetime64[ns](1), float64(1), int64(4), object(6)
memory usage: 167.1+ KB
```

1.5 Store

```
In [234]: # store the final dataframe "df" in the csv file.
df.to_csv("twitter_archieve_master.csv", index=False)
```

Exploratory Data Analysis

```
In [235]: pd.plotting.scatter_matrix(df.iloc[:,5:9], figsize=(8,8));
```



1.5.1 Topic1. How the ratings are related with stages and types of dogs

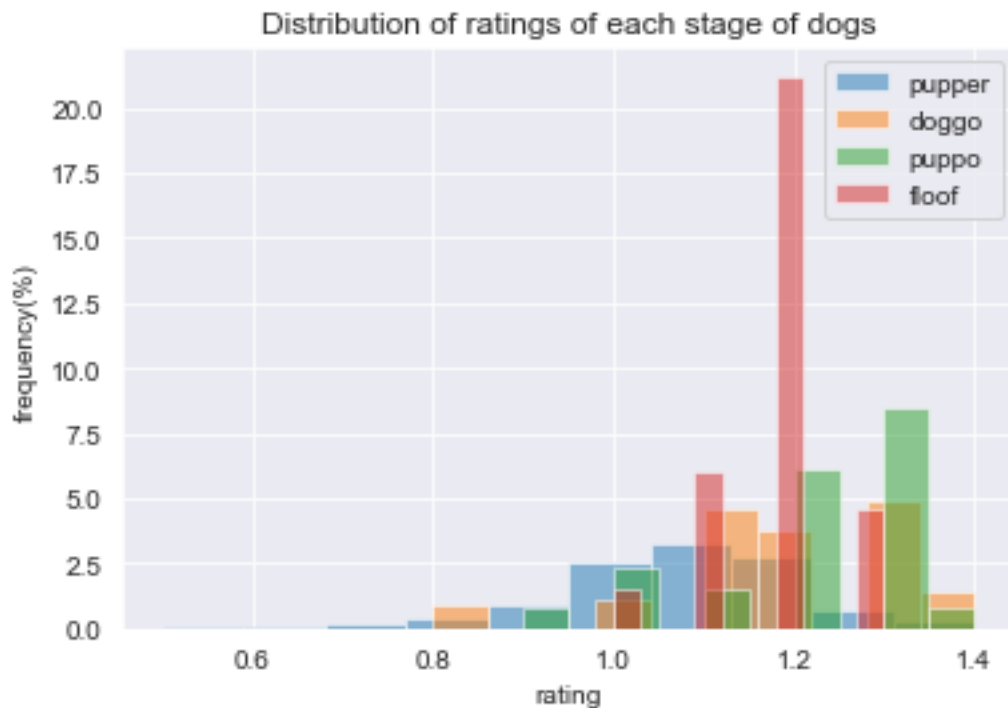
It can be possible that a master of the account has personal preference of dog stages or dog types and the rating is somehow biased. We take a look at the relationship between dog ratings, stages, and types.

First, we visualize the relationship between dog stages and ratings. The dog stage "blep" has only 1 tweet, so we ignore this stage.

```
In [236]: sns.set_style("darkgrid")

# visualize each stage's distribution
df[df.stage == "pupper"]["rating"].hist(density=1, alpha=0.5, label="pupper")
df[df.stage == "doggo"]["rating"].hist(density=1, alpha=0.5, label="doggo")
df[df.stage == "puppo"]["rating"].hist(density=1, alpha=0.5, label="puppo")
df[df.stage == "floof"]["rating"].hist(density=1, alpha=0.5, label="floof")
plt.title("Distribution of ratings of each stage of dogs")
```

```
plt.xlabel("rating")
plt.ylabel("frequency(%)")
plt.legend();
```



What we have found is that there seems to be a little effect on the ratings depending on the stage of dogs. The graph suggests that the stage “puppo” tend to have the highest rating, followed by “floof”, “doggo”, and “pupper”.


Next, we visualize the relationship between dog types and ratings. We begin by knowing dog type category names.

```
In [237]: # get to know the dog type categories
df.dog_type.value_counts().sort_values(ascending=False).head(10)
```

```
Out[237]: golden_retriever    152
Labrador_retriever    106
Pembroke                93
Chihuahua              87
pug                    62
toy_poodle             50
chow                   48
Samoyed                42
Pomeranian            41
malamute               33
Name: dog_type, dtype: int64
```

There are too many types of dogs, so we focus only on the 5 most types which are “golden_retriever”, “Labrador_retriever”, “Pembroke”, “Chihuahua”, and “pug”.

```
In [238]: # visualize each type's distribution
df[df.dog_type == "golden_retriever"]["rating"].hist(density=1, alpha=0.5, label="golden_retriever")
df[df.dog_type == "Labrador_retriever"]["rating"].hist(density=1, alpha=0.5, label="Labrador_retriever")
df[df.dog_type == "Pembroke"]["rating"].hist(density=1, alpha=0.5, label="Pembroke")
df[df.dog_type == "Chihuahua"]["rating"].hist(density=1, alpha=0.5, label="Chihuahua")
df[df.dog_type == "pug"]["rating"].hist(density=1, alpha=0.5, label="pug")
plt.title("Distribution of ratings of each type of dogs")
plt.xlabel("rating")
plt.ylabel("frequency(%)")
plt.legend();
```



Apparent gaps have not found between types, but “pug” seems to tend to get rather low ratings.

Now we use multiple linear model for the relationship between ratings and the combination of stages and columns. (In this project, let me skip to check if this really has no problem in terms of multicollinearity .etc.) We set “pupper” and “pug” as base-lines and see if the differences of ratings from other each stage or types are statistically significant or not.

```
In [239]: # extract rows which are given both specific stage and type
df2 = df.copy()
```

```

df2 = df2[df2.stage.isin(["puppo", "pupper", "doggo", "floof"])]
df2 = df2[df2.dog_type.isin(["golden_retriever", "Labrador_retriever", "Pembroke", "Chihuahua"])]

# add intercept and dummy columns of dog stages and types
df2["intercept"] = 1.
df2[["doggo", "floof", "pupper", "puppo"]] = pd.get_dummies(df2.stage)
df2[["Chihuahua", "Labrador_retriever", "Pembroke", "golden_retriever", "pug"]] = pd.get_dummies(df2.dog_type)
df2.sample(3)

```

```

Out[239]:
      tweet_id      timestamp      source \
0  890240255349198849  2017-07-26 15:59:51  <a href="http://twitter.co...
458 693262851218264065  2016-01-30 02:41:58  <a href="http://twitter.co...
445 699434518667751424  2016-02-16 03:25:58  <a href="http://twitter.co...

      text      expanded_urls  rating \
0  This is Cassie. She is a c...  https://twitter.com/dog_ra...  1.4
458 I hope you guys enjoy this...  https://twitter.com/dog_ra...  1.1
445 I know this is a tad late ...  https://twitter.com/dog_ra...  1.2

      stage  ret  fav      jpg_url ...  intercept  doggo \
0  doggo  7300  31555  https://pbs.twimg.com/medi...  1.0  1
458 pupper  539  2358  https://pbs.twimg.com/medi...  1.0  0
445 pupper  547  2287  https://pbs.twimg.com/medi...  1.0  0

      floof  pupper  puppo  Chihuahua  Labrador_retriever  Pembroke \
0  0  0  0  0  0  1
458 0  1  0  0  0  0
445 0  1  0  0  0  0

      golden_retriever  pug
0  0  0
458 1  0
445 1  0

```

[3 rows x 22 columns]

```

In [240]: # apply multiple linear regression model
lm = sm.OLS(df2.rating, df2[["intercept", "doggo", "floof", "puppo", "Chihuahua", "Labrador_retriever", "Pembroke", "golden_retriever", "pug"]])
results = lm.fit()
results.summary()

```

```

Out[240]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
Dep. Variable:                  rating    R-squared:                0.274
Model:                            OLS    Adj. R-squared:            0.212
Method:                    Least Squares    F-statistic:                4.425

```

```

Date:          Fri, 12 Oct 2018    Prob (F-statistic):      0.000326
Time:          01:22:52           Log-Likelihood:         75.298
No. Observations:      90          AIC:              -134.6
Df Residuals:          82          BIC:              -114.6
Df Model:              7
Covariance Type:      nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975
-----
intercept          1.0182      0.033     30.753      0.000      0.952      1.084
doggo              0.0900      0.029      3.053      0.003      0.031      0.149
floof             0.0287      0.049      0.589      0.558     -0.068      0.125
puppo             0.0622      0.045      1.371      0.174     -0.028      0.152
Chihuahua         0.1174      0.050      2.358      0.021      0.018      0.217
Labrador_retriever 0.1015      0.044      2.329      0.022      0.015      0.188
Pembroke          0.1618      0.045      3.631      0.000      0.073      0.250
golden_retriever   0.1237      0.040      3.098      0.003      0.044      0.203
=====
Omnibus:              4.612    Durbin-Watson:          1.713
Prob(Omnibus):        0.100    Jarque-Bera (JB):        3.882
Skew:                 -0.463    Prob(JB):                0.144
Kurtosis:              3.423    Cond. No.                8.12
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

This results suggests that (in terms of dog stages) “doggo” has statistically significant difference from “pupper” and that (in terms of dog types) each of “Chihuahua”, “Labrador_retriever”, “Pembroke”, “golden_retriever” has statistically significant difference from “pug”.

Also, it can be said that 27.4% of rating evaluation can be “explained” by these 2 factors.

1.5.2 Answer1. Relationship between stages of dogs and the ratings.


- Regarding dog stages, it is statistically significant that “doggo” tend to get better ratings than “pupper”.
- Regarding dog types, it is statistically significant that “Chihuahua”, “Labrador_retriever”, “Pembroke”, and “golden_retriever” tend to get better ratings than “pupper”.
- 27.4% of rating evaluation juege depends on these 2 factors.

1.5.3 Topic2. Relationship between favorite counts and the ratings

In this data, we have 2 indicators that show reader’s preference on tweets, which are “favorite” counts and “retweet” counts. It is easily predictable that tweets with high

ratings are favored and also retweeted more and also that there is a strong positive correlation between “favorite” counts and “retweet” counts. First, this should be checked.


```
In [241]: plt.scatter(df.fav, df.ret)
plt.title("Relationship between favorite-count and retweet-count" )
plt.xlabel("Favorite count")
plt.ylabel("Retweet count");
```



As we predicted, these counts are correlated. This suggests that we can deem favorite counts as a reliable indicator of tweets’ popularity. In this analysis, we focus on “favorite” counts. (And also I have noticed that there is a tweet which had by far the most favs and retweets from the graph. This would be assessed visually later.)

We take a look at the relationship between ratings from the account and the favorite counts.

```
In [242]: plt.scatter(df.rating, df.fav)
plt.title("Relationship between ratings and favorite-count")
plt.ylabel("Favorite count")
plt.xlabel("Rating");
```



We use simple linear regression model to interpret this result (even though it may not be the best way.)

```
In [243]: df["intercept"] = 1.
lm = sm.OLS(df.fav, df[["intercept","rating"]])
results = lm.fit()
results.summary()
```

```
Out[243]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  fav    R-squared:                0.178
Model:                            OLS    Adj. R-squared:           0.177
Method:                 Least Squares    F-statistic:              355.2
Date:                  Fri, 12 Oct 2018    Prob (F-statistic):       7.00e-72
Time:                   01:22:55    Log-Likelihood:          -17804.
No. Observations:                1645    AIC:                     3.561e+04
Df Residuals:                    1643    BIC:                     3.562e+04
Df Model:                        1
Covariance Type:                nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----

```

intercept	-2.589e+04	1885.466	-13.729	0.000	-2.96e+04	-2.22e+04
rating	3.236e+04	1716.991	18.846	0.000	2.9e+04	3.57e+04
=====						
Omnibus:		1648.536	Durbin-Watson:			1.691
Prob(Omnibus):		0.000	Jarque-Bera (JB):			101955.703
Skew:		4.749	Prob(JB):			0.00
Kurtosis:		40.381	Cond. No.			12.6
=====						

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified
"""
```

Even though this suggests that there is a positive correlation between rating result and favorite accounts, R-squared value(0.178) tells us that there may be a better way to understand this relationship.

1.5.4 The best tweet which is extremely popular

We take a look at the tweet which had extremely popularity

```
In [244]: best_tweet = df[df.fav == df.fav.max()]
          best_tweet
```

```
Out[244]:
```

	tweet_id	timestamp	source	\
139	744234799360020481	2016-06-18 18:26:18	<a href="http://twitter.co...	

	text	expanded_urls	rating	\
139	Here's a doggo realizing y...	https://twitter.com/dog_ra...	1.3	

	stage	ret	fav	jpg_url	img_num	\
139	doggo	84426	165057	https://pbs.twimg.com/ext...	1	

	dog_type	intercept
139	Labrador_retriever	1.0

```
In [245]: pd.options.display.max_colwidth = 150
          best_tweet.text
```

```
Out[245]: 139    Here's a doggo realizing you can stand in a pool. 13/10 enlightened af (vid by
          Name: text, dtype: object
```

I checked the tweet visually by jumping to the link, but what I have found was just a movie in which a dog is swimming in the pool. It may be said that an ordinary tweet sometimes attracts enormous attention unexpectedly.

1.5.5 Answer2. Relationship between favorite counts and the ratings

- There is a positive correlation between rating result and favorite accounts.

- This relationship is not completely explained by a linear model.

Conclusions

In this analysis, we focused on the tweet data from a dog-rating twitter account. We began from getting data from api. Then we cleaned and finally analyzed the data.

We found that some dog stages and types get higher ratings from the account. Favorite count, which is a great indicator of the popularity of tweets, tend to be more for the tweets which are more highly evaluated. However, it is not sure whether reader favored because they found that the rating better, or because pictures themselves are wonderful worth great rating.

Reference

- tweet data of “WeRateDogs (@dog_rates)” from https://twitter.com/dog_rates
- ‘image_predictions.tsv’ file from “<https://d17h27t6h515a5.cloudfront.net/topher/2017/August/599fd2ad-predictions/image-predictions.tsv>”